

Runner.ai memory architecture design.

The memory system uses rolling thread summaries plus recent messages.

Long chats are compacted into a summary block to fit LLM context.

Every tool observation is returned to the LLM as a ToolMessage.

Retrieved evidence is normalized centrally before being passed onward.